



Paused training | Global batch 256 · Replay microbatch 32/GPU · 8 GPUs

OPD: policy-gradient objective



Stage3: forward KL + 0.1 × attention loss



Raw per-update losses; changing batches and different objectives — not a quality comparison.