



Stage2 C256 / GB128 / 8 × A100 — training and held-out loss

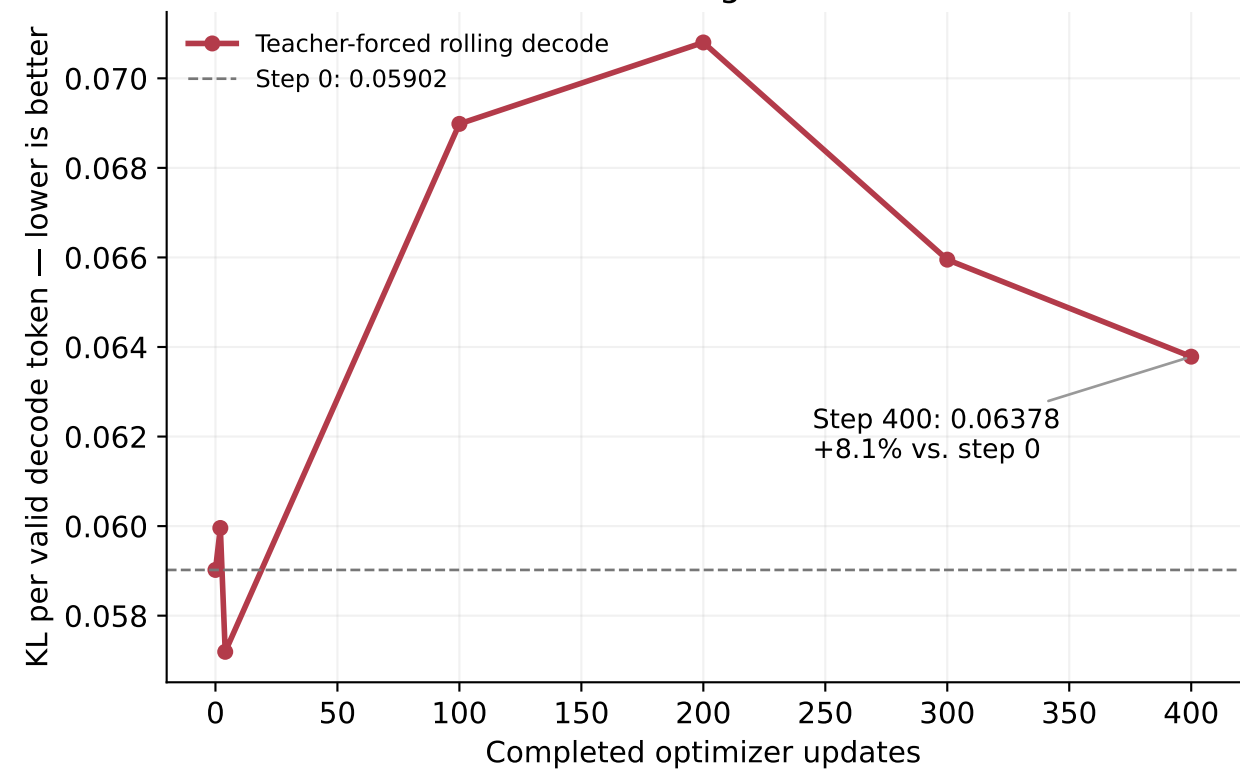
Training objective: KL + 0.1 × attention auxiliary loss



Held-out prefill KL vs. step 0



Held-out rolling decode KL



Validation: fixed 16 examples; 22,908 prefill and 20,706 decode positions. Resume duplicates removed.
Step 0 = completed Stage1 weights. KL diagnostics do not measure free-generation task accuracy.